

# FROM TRACES TO GUARDED PROGRAMS: EVIDENCE-GATED COMPILATION OF RECURRENT AGENT WORKFLOWS

**Anonymous authors**

Paper under double-blind review

## ABSTRACT

Tool-using agents often revisit the same read-only evidence path while paying for a model decision at every tool boundary. Replacing a recurring model-mediated prefix could save these calls, but it is not ordinary caching: a later tool argument may encode a decision unrecoverable from observed state, an apparently read-only call may cross an effect boundary, and correct tool replay may still change the remaining agent’s answer. Existing recurrence, replay, and workflow-induction signals therefore identify candidate regions, but do not establish that replacing the model is safe.

We introduce guarded agentic compaction (GAC), a compile-or-retire trace-to-program system that makes this evidence requirement explicit. GAC normalizes recorded episodes, mines recurrent read-only prefixes, and emits a bounded deterministic program only when every tool argument has a provenance witness from entry state or a prior result; an unwitnessed argument retires the region while preserving its maximal justified prefix. Manifest, input, and output guards enforce the resulting contract, and a fixed finite-sample selective-risk test decides deployment; a failed check falls back to the unchanged agent. We evaluate this lifecycle on three live GitHub workflow families, a time-forward five-repository extension, and four public compiler substrates. On the live families, GAC attains 90/90 exact held-out contracts, versus 89/90 for the unchanged agent, while reducing provider requests by 66.6% and tokens by 63.1% in aggregate. It completes four of five extension repositories and retires one. Three substrates retire every recurrent family despite successful provenance, synthesis, and replay; the fourth is the first whose sample size makes admission reachable, and it admits—then 8,190 released official-baseline trajectories show the admitted region is structurally eligible on 2,339/2,340 full-code trajectories but only 1/4,680 ReAct or plan-and-execute trajectories, so deployability is a precondition separate from admissibility. GAC thus turns recurrence from permission to rewrite an agent into a testable, evidence-gated specialization decision with principled refusal and fallback.

## 1 INTRODUCTION

The standard agent loop pays for a model decision at every tool boundary: model, tool, model, tool, and so on (Yao et al., 2023). That flexibility is valuable while a workflow is novel. At scale, however, mature agents often revisit the same read-only evidence path with the same data dependencies. In those settings, repeating the model decision adds latency, tokens, and cost without necessarily adding useful adaptation.

Removing that decision is not ordinary caching. A tool argument may depend on a specific earlier observation, a repeated sequence may cross an approval or write, an apparently read-only tool may hide an undeclared effect, and a program that replays the same tool outputs may still change the downstream answer. The problem is therefore not just “find recurring traces.” The problem is to decide when a recurrent model-mediated region can be replaced by a deterministic program and when it should be refused.

Consider held-out issue #4420 in the expanded 30-record study. Given only `issue_number=4420`, the observed trace calls `record(4420)` and then passes the `returned` `record.issue_number` to `labels` and to `comments`, which it invokes with `limit=100` before rendering an auditable triage




Figure 1: Two retained cohorts expose different failure modes. In the expanded study, issue #4420 forces a shorter grounded prefix; in the earlier study, issue #6602 shows that clean tool replay need not preserve the downstream answer.

answer. All 132 discovery executions take that route, so a recurrence-only optimizer would emit all three reads. GAC instead retires the *region*: `comments.limit` varies across otherwise similar entries and is reachable neither from the entry state nor from either earlier result, so one unwitnessed slot blocks the whole three-read candidate as `ungroundable_slot`. What survives is the maximal justified prefix `record → labels`; the agent still chooses the comment limit and writes the answer. That two-read artifact passes 92/92 calibration groups at  $\alpha = .05$  (upper bound 0.0498).

Why stop there? A separate, earlier 18-record study supplies the counterexample: its aggressive three-read artifact replayed all 45 tool calls, yet issue #6602 lost the URL from a Markdown link and only 17/18 downstream answers remained exact. A provider-free continuation guard detects and checked-renders that miss; it is a mechanism check, not live safety evidence. Together the two cohorts (fig. 1) show why recurrence proposes a candidate, while provenance, continuation scope, and statistical evidence determine what may replace the model.

Framed this way, the closest precedent is not prompt optimization but the tracing just-in-time compiler. Dynamo (Bala et al., 2000) detects a hot path, compiles it behind a *guard*, and deoptimizes to the interpreter when the guard fails (Hölzle et al., 1992). Agent traces admit the same structure — hot-trace detection becomes family mining, guard synthesis becomes an entry contract, and deoptimization becomes fallback to the unchanged agent — but they add three obligations a JIT does not face. Tool arguments must be shown to be *grounded* in observable state rather than invented by the model; declared *effects* must permit speculation, because an agent’s “interpreter” can call the outside world; and the residual error rate of a compiled region is statistical, so it needs a finite-sample bound rather than a soundness proof.

We study these obligations through *guarded agentic compaction* (GAC), a compile-or-retire trace-to-program system for recurrent read-only prefixes (fig. 2). The core contribution is not “agent compilation” in general, but an admissibility argument for trace-derived specialization:

1. A guarded specialization pipeline (Alg. 1) that combines typed provenance, effect and position barriers, bounded synthesis, runtime verification, and finite-sample compile-or-retire admission (Alg. 2, proposition 1); every stage can refuse, and the default output is the unchanged agent.
2. A real-provider evaluation across three public-record workflow families, plus a narrower cross-repository, time-forward extension that keeps both successful artifacts and principled retirements.
3. A negative-result boundary showing that recurrence and replayability do not themselves license deployment: NESTFUL, API-Bank, and executed BFCL gold plans retire under the same admission logic, and hand-written programs are competitive or tied wherever the workflow is already known, so no claim of automatic superiority is warranted.
4. A pre-registered fourth substrate, AppWorld, which is the first external corpus whose sample size makes admission reachable and which does admit — and a measurement on 8,190 released official-baseline trajectories showing that the admitted region is structurally eligible on 2,339/2,340 full-code trajectories but only 1/4,680 ReAct or plan-and-execute trajectories, which separates *admissible* from *dispatchable*.

## 2 PROBLEM FORMULATION

**Episodes and candidate regions.** An episode  $E = (z, M, e_{1:T}, y)$  contains an entry-state snapshot  $z$ , a content-addressed execution manifest  $M$ , ordered events  $e_t$  (model requests and responses, tool calls and results, handoffs, guardrails, approvals, errors, commit boundaries), and an observable outcome  $y$ . Each tool  $f$  carries an application-owned effect declaration  $\epsilon(f)$  and capability flags such as *speculatable*; an unknown effect is not a read.

A candidate region  $R = [a, b]$  begins at a model-request boundary and ends after a tool result. It is *groundable* when every call argument is either an application-declared literal for that slot or an expression over entry state or prior results inside  $R$ ,

$$\forall c_j \in R, \forall u \in \text{args}(c_j) : \text{Lit}(c_j, u) \vee [u = g_{j,u}(z, o_{<j}), g_{j,u} \in \mathcal{L}], \quad (1)$$

where  $\text{Lit}$  is a schema-level allowlist fixed independently of the traces and  $\mathcal{L}$  is a closed, bounded transform library; and *effect-admissible* when every tool in  $R$  is declared read-only, speculatable, replayable, and inside one principal and isolation partition. Handoffs, approvals, writes, errors, and unknown effects are hard barriers, not costs to be traded away. Because the deployed runtime resolves artifacts at the initial model boundary, the compiler also imposes a *position invariant*:

$$a = 0 \quad \text{in the normalized model-boundary index.} \quad (2)$$

This is not aesthetic. A suffix program dispatched at entry can reorder or duplicate calls the agent has already made — a semantic mismatch between compilation-time and resolution-time position rather than a tuning failure, and a pilot that violated eq. (2) produced exactly that fault (appendix E).

**Selective optimization objective.** The compiler emits an artifact  $A = (P, H, V, q, \eta, M)$ : a deterministic program  $P$ , hard guard  $H$ , output verifier  $V$ , nonconformity score  $q$ , threshold  $\eta$ , and manifest pins  $M$ . Let  $x = (z, M', \omega)$  denote a future entry state, runtime manifest, and external state. Here  $M' \simeq M$  means equality on every pinned compatibility field;  $\omega$  can affect outcomes and cost but is not observable at the dispatch boundary. Dispatch is *selective*,

$$d_A(x) = \mathbf{1}\{H(z, M') = 1 \wedge q(z) \leq \eta \wedge M' \simeq M\}, \quad (3)$$

and any failed precondition runs the unchanged baseline agent  $B$ . If execution fails before a commit and staging proves the attempt clean,  $B$  runs; a dirty failure after a commit is an incident, not a fallback. Fix  $L(A, x) = 1$  for an externally visible task/contract violation or incident after dispatch; a clean abort followed by unchanged  $B$  has  $L = 0$ . Let  $C(A, x)$  include the complete realized cost of that attempt, including a clean fallback. For a future group  $G = (x_1, \dots, x_m)$ , define  $D_A(G) = \max_i d_A(x_i)$  and  $W_A(G) = \max_i d_A(x_i)L(A, x_i)$ . We seek

$$\max_A \mathbb{E}[d_A(X) \{C(B, X) - C(A, X)\}] \quad \text{s.t.} \quad \Pr[W_A(G) = 1 \mid D_A(G) = 1] \leq \alpha. \quad (4)$$

Two properties shape everything that follows. The constraint conditions on dispatch: abstention does not enter its denominator, although it earns no savings in the objective; by convention its conditional risk is zero when population dispatch probability is zero. Returning *no* artifact is feasible, so RETIRE




Figure 2: Architecture of GAC. (A) Capture normalizes traces into one typed Episode IR against a declared effect catalog. (B) Offline compilation is a cascade of independent rejection points; the shaded band lists barriers no statistical evidence can override. (C) At runtime exactly one terminal edge compacts; the rest return the unchanged agent.

is a legitimate output. Equation (4) is a design target: the implementation calibrates a bounded ranked set and retains nondominated survivors, and we prove neither optimality nor a compiler-wide guarantee for that search. Calibration uses independent scenario groups as its statistical units. The group and episode risks coincide in the reported live studies, where every group contains exactly one public record; they need not coincide for unequal multi-episode groups.

### 3 GUARDED AGENTIC COMPACTION

Figure 2 and Alg. 1 give the same six stages at two levels of detail. The order is itself part of the argument: each stage can only weaken a candidate, so a later stage may reject one an earlier stage passed, but no later stage can authorize one that an earlier stage blocked. We describe each and state what it can refuse.

**1. Typed provenance.** A capture adapter normalizes framework traces into the Episode IR; manifests pin prompt, policy, guardrail, tool, model, SDK, tracer, entry-contract, and effect-catalog identities, and episodes with different compatibility keys are partitioned before any learning. For each call-argument slot, the program-argument trace graph (PATG, Alg. 3) searches entry-state and prior-result paths for typed values reachable by a bounded transform — execution provenance in the database sense (Cheney et al., 2009), specialized to tool arguments. Two outcomes block a region rather than guessing: a slot with *no* witness is a genuine model decision, and a slot with more than  $\kappa = 3$  witnesses is ambiguous. Figure 3 shows both branches on the motivating record. A repeated tool name, a matching literal, and a successful replay are not witnesses; only an expression in  $\mathcal{L}$  is.

**2. Region mining and ranking.** The miner enumerates model-boundary-to-result windows of at most  $w_{\max}$  tool calls, qualifies live-ins and effects against a human-signed rather than inferred catalog, and hashes tool signatures, dependency topology, run shape, and live-in/live-out shape while abstracting literal values. Families require support across independent scenario *groups* rather than raw event

**Algorithm 1** COMPILER — the compile-or-retire cascade. Every candidate is either registered or assigned an attributed RETIRE, so refusal accounting remains explicit. Stages marked (CALLER) are separate library entry points run in this order rather than steps inside one function.

**Input:** episodes  $\mathcal{E}$ ; effect catalog  $C$ ; entry schema  $\Theta$ ; transform library  $\mathcal{L}$ ; ambiguity cap  $\kappa$ ; budgets  $\alpha, \delta$ ; grid  $\Lambda$ ; feasibility target  $\Delta$

**Output:** artifact set  $\mathcal{A}$ , or RETIRE with attributed reasons

```

1:  $\mathcal{A} \leftarrow \emptyset$ 
2:  $\mathcal{E} \leftarrow \text{QUALIFY}(\mathcal{E}, C)$ ;  $\{\mathcal{E}_M\} \leftarrow \text{PARTITIONBY}(\mathcal{E}, \text{manifest}, \text{principal}, \text{isolation})$   $\triangleright$  drop partial runs,
   uninned manifests
3: for all partitions  $\mathcal{E}_M$  do
4:    $(\mathcal{T}, \mathcal{D}, \mathcal{K}, \mathcal{S}) \leftarrow \text{GROUPEDSPLIT}(\mathcal{E}_M)$  (CALLER)  $\triangleright$  train / dev / calibration / sealed test, split by group
5:    $G \leftarrow \text{BUILDPATG}(\mathcal{T}, \Theta, C, \mathcal{L}, \kappa)$   $\triangleright$  Alg. 3: block ungrounded or ambiguous slots
6:    $\mathcal{F} \leftarrow \text{CANONICALIZE}(\text{MINEREGIONS}(G, C, w_{\max}))$ , keeping only cross-group support
7:   if  $\text{CEILING}(\mathcal{T}, \mathcal{F}) < \Delta$  then (CALLER)
8:     record RETIRE (infeasible ceiling, Eq. 6); continue  $\triangleright$  try the next compatible partition
9:   end if
10:  for all families  $F \in \mathcal{F}$  ranked by Eq. (5) do
11:     $P \leftarrow \text{SYNTHESIZE}(F, \mathcal{L})$  at depth  $\leq 2$ ; if  $P = \perp$  then record RETIRE (synthesis); continue
12:     $(H, V) \leftarrow \text{INDUCECONTRACT}(F)$ ; if  $\neg \text{REPLAYANDCHALLENGE}(P, V, \mathcal{D})$  then record RETIRE
      (counterexample); continue
13:     $q \leftarrow \text{FITSCORE}(\mathcal{D})$ ; freeze  $q$   $\triangleright$  dev groups only; never calibration
14:     $\eta \leftarrow \text{CALIBRATE}(P, H, V, q, M, \mathcal{K}, \Lambda, \alpha, \delta)$ ; if  $\eta = \perp$  then record RETIRE (risk); continue  $\triangleright$ 
      Alg. 2
15:     $\mathcal{A} \leftarrow \mathcal{A} \cup \{\text{PACKAGE}(P, H, V, q, \eta, M)\}$  with evidence and lifecycle
16:  end for
17: end for
18: return  $\text{SELECTUNDOMINATED}(\mathcal{A})$  if nonempty, else RETIRE

```




Figure 3: Equation (1) on the held-out issue-4420 record. Every argument slot needs an expression in  $\mathcal{L}$  over the entry state or a prior result; the two `issue_number` slots have one and the `limit` slot has none, so it is drawn as a socket with no incoming edge. The question is strictly harder than recurrence: not “did this sequence repeat?” — all 132 discovery executions did — but “can every socket be wired from an allowed prior value?” One unwitnessed slot retires the whole region, and a slot with more than  $\kappa$  witnesses does the same.

frequency, which prevents a single chatty scenario from manufacturing recurrence. Survivors are ranked by an MDL-inspired trade-off that pays for expected savings and charges for heterogeneity, effect exposure, and program size:

$$\text{score}(F) = \underbrace{s_F \bar{k}_F c_m}_{\text{support-weighted saving}} - \lambda_1 \text{Ent}(F) - \lambda_2 \rho_{\text{eff}}(F) - \lambda_3 |F|, \quad (5)$$

with  $s_F$  the number of supporting train groups,  $\bar{k}_F$  the mean removable model requests,  $c_m$  their unit cost, and  $\text{Ent}(F)$  the Shannon entropy over canonical variants inside the family; without  $\lambda_1$  the ranking prefers a heterogeneous family whose single canonical form fits none of its members. Equation (5) affects only *which* family is tried first — no admission decision depends on it.

**Algorithm 2** CALIBRATE — fixed-grid exact selective admission. The score  $q$  and the grid  $\Lambda$  are frozen *before* this procedure sees a calibration group, which is what makes the union bound legitimate.

**Input:** frozen candidate  $A_0 = (P, H, V, q, M)$  without threshold; calibration groups  $\mathcal{K}$  with fixed per-instance violations  $L(A_0, x)$ ; grid  $\Lambda$ ; budgets  $\alpha, \delta$

**Output:** threshold  $\eta$  with a simultaneous guarantee, or  $\perp$  (RETIRE)

```

1:  $\mathcal{A} \leftarrow \emptyset$  ▷ admissible (coverage, threshold) pairs
2: for all  $\eta \in \Lambda$  do
3:   define  $d^\eta(x)$  by Eq. (3) with threshold  $\eta$ ;  $D^\eta(g) \leftarrow \max_{x \in g} d^\eta(x)$ ;  $W^\eta(g) \leftarrow \max_{x \in g} d^\eta(x)L(A_0, x)$ 
4:    $K_\eta \leftarrow \{g \in \mathcal{K} : D^\eta(g) = 1\}$ ;  $n_\eta \leftarrow |K_\eta|$ ; if  $n_\eta = 0$  then continue ▷ vacuous bound  $U_\eta = 1$ 
5:    $k_\eta \leftarrow |\{g \in K_\eta : W^\eta(g) = 1\}|$  ▷ wrong after dispatch only, never abstention
6:    $U_\eta \leftarrow 1$  if  $k_\eta = n_\eta$ , else  $\text{Beta}^{-1}(1 - \frac{\delta}{|\Lambda|}; k_\eta + 1, n_\eta - k_\eta)$ ; if  $U_\eta \leq \alpha$  then  $\mathcal{A} \leftarrow \mathcal{A} \cup \{(n_\eta/|\mathcal{K}|, \eta)\}$ 
▷ Eq. (7)
7: end for
8: return  $\perp$  if  $\mathcal{A} = \emptyset$ , else the  $\eta$  component of  $\text{lexmax}_{(\text{cov}, \eta)} \mathcal{A}(\text{cov}, \eta)$  ▷ largest empirical group coverage; larger threshold breaks ties

```

**3. Feasibility ceiling.** Before any synthesis runs, an estimator bounds what the compiler could achieve with a perfect verifier and a gate that never abstains:

$$\Delta_{\max} = \frac{\varphi k}{n_B}, \quad \text{necessary: } \varphi \rho k \geq \Delta n_B, \quad (6)$$

with  $n_B$  baseline model requests per episode,  $\varphi$  the fraction of episodes holding an eligible region,  $k$  the requests removed per successful dispatch, and  $\rho$  the verifier pass rate. Since eq. (6) assumes  $\rho = 1$ , no measured reduction can exceed it, and it is met with equality exactly when nothing abstains and nothing fails. Reporting it first makes a decisive negative answer cheap: a workload whose ceiling sits under the target can be declined without building a compiler for it.

**4–5. Bounded synthesis and contracts.** Synthesis searches a 23-operator library at depth at most two, covering path selection, indexing, string normalization, arithmetic, comparisons, bounded collection operations, and narrow loops; bindings are ranked by stability across groups, and a group-wise refit rejects bindings that exploit row-level coincidence. The hypothesis space is deliberately small enough to read: the approach resembles programming-by-example (Gulwani, 2011) and bounded sketching (Solar-Lezama et al., 2006), and it does *not* generate arbitrary Python. Hard guards then constrain entry types, required fields, categorical sets, numeric intervals, patterns, isolation keys, allowed effects, and manifest pins; verifiers constrain live-out presence, type, cardinality, empirical hulls, provenance, effect multisets, and call counts. These are candidate invariants, not proofs for all future inputs — which is why a statistical gate follows them rather than replacing them.

**6. Exact risk-gated admission.** Algorithm 2 decides deployment. The score  $q$  is fitted on *dev* groups only — never on calibration groups — from entry-observable features such as missing fields, hull margin, temporal drift, and provenance ambiguity, then frozen together with the finite threshold grid  $\Lambda = \{.02, .05, .08, .11, .14, .17, .20, .25, .30, .40, .50\}$ . Two data signals do two different jobs here, and conflating them is the easiest way to misread the guarantee:  $q$  is *fitted* on development groups that turned out unproductive in any way — wrong or abstained — while the bound below is *counted* only from groups that were dispatched and turned out wrong. For each  $\eta \in \Lambda$  the compiler counts admitted calibration groups ( $n_\eta$ ) and those containing at least one post-dispatch violation ( $k_\eta$ ), and inverts the one-sided exact binomial test (Clopper & Pearson, 1934) at a Bonferroni-corrected level, with explicit empty and all-violation cases:

$$U_\eta = \begin{cases} 1, & n_\eta = 0 \text{ or } k_\eta = n_\eta, \\ \text{Beta}^{-1}\left(1 - \frac{\delta}{|\Lambda|}; k_\eta + 1, n_\eta - k_\eta\right), & 0 \leq k_\eta < n_\eta, \end{cases} \quad U_\eta|_{k_\eta=0} = 1 - \left(\frac{\delta}{|\Lambda|}\right)^{1/n_\eta}. \quad (7)$$

It admits the largest-coverage threshold with  $U_\eta \leq \alpha$  and retires the family otherwise — a fixed-grid instance of the learn-then-test discipline (Angelopoulos et al., 2022). The closed form on the right is the operative constraint in our experiments: at  $\alpha = 0.05$ ,  $\delta = 0.1$ , and  $|\Lambda| = 11$ , a zero-violation family still needs  $n_\eta \geq 92$  admitted groups. The default output is refusal, not emission.

**Proposition 1** (Per-candidate selective-risk admission). *Condition on train/development data and fix one candidate  $(P, H, V, q)$ , group-level violation rule, and finite grid  $\Lambda$  before observing calibration. If calibration groups are i.i.d. from the future-group distribution, then with probability at least  $1 - \delta$ , every  $\eta \in \Lambda$  satisfies  $r_\eta \leq U_\eta$ , where  $r_\eta$  is  $\Pr(W^\eta = 1 \mid D^\eta = 1)$  when population admission is positive and is defined as zero otherwise. Consequently the selected  $\hat{\eta}$  satisfies  $r_{\hat{\eta}} \leq \alpha$  whenever  $U_{\hat{\eta}} \leq \alpha$ .*

Appendix A gives the proof. The guarantee is deliberately narrow: it is *per fixed candidate*, not compiler-wide over the adaptive candidate search Alg. 1 actually performs, and it conditions on group exchangeability rather than establishing it (section 7). The scientific point is that guarded deployment requires an explicit evidence boundary, not just a recurrent trace and a successful replay.

## 4 EXPERIMENTAL SETUP

Because the claim is an *admissibility* claim, the evaluation asks three questions: does specialization preserve outcomes while saving work on real records (section 5.1)? Does it transfer under a time-forward split (section 5.2)? And does it refuse when evidence is thin (section 5.3)?

### 4.1 BENCHMARK AT A GLANCE

Every GitHub sample begins with one issue or pull-request number from a pinned public snapshot. The agent must make read-only calls and return exact, source-grounded fields, not a subjective preference. The three primary families use 132 discovery and 30 balanced held-out records each, with live provider calls. Each family pairs a read-only evidence path with an exact answer contract — issue type, pull-request outcome, or backlog attention route, each returning a category plus title and a verbatim source excerpt — so the benchmark tests both the structured answer and whether its evidence survives compaction (table 3).

**Cross-repository time-forward extension.** A narrower two-read PR-outcome task spans five frozen public repositories. Held-out records are strictly newer than discovery records, and selection is sealed before provider calls; repositories without an admitted artifact remain in the result. A balanced rerun on three repositories controls the open/merged/closed\_unmerged mix.

**External compiler substrates.** NESTFUL, API-Bank, and executed BFCL v4 gold plans supply valid executable traces with the argument-level producer structure provenance, synthesis, replay, and admission need, and measure none of them as live planning quality (Basu et al., 2025; Li et al., 2023; Patil et al., 2025). AppWorld adds a fourth substrate and a different one (Trivedi et al., 2024): its public gold solutions, executed on its pinned resettable backend, give 136 byte-exactly reproducible traces over 68 typed APIs and enough groups for eq. (7) to be satisfiable, and its released baseline runs supply held-out trajectories for the dispatch measurement. No model runs on any substrate, so none licenses an accuracy, quality, or efficiency claim, and none is pooled with the GitHub families (appendix F.3).

**Conditions, metrics, and statistics.** Each GitHub family compares the unchanged baseline  $B$ , the compiled artifact, and a hand-written pre-model program on the same records, so the comparison is paired rather than across cohorts. Quality is exact task-contract satisfaction; resources are requests, interfaces, tokens, latency, and estimated cost. We use paired exact McNemar tests, bootstrap intervals, and Wilcoxon signed-rank tests; the only guaranteed multiplicity control is the Bonferroni term in eq. (7). All admission runs use  $\alpha = 0.05$ ,  $\delta = 0.1$ , and  $|\Lambda| = 11$  (appendix C).

## 5 RESULTS

### 5.1 PRIMARY WORKFLOWS: PRESERVED CONTRACTS, LESS WORK

Across the 90 held-out records, GAC satisfies 90/90 exact contracts versus 89/90 for the unchanged agent (table 1), reducing provider requests by 66.6%, tokens by 63.1%, latency by 64.2%, and estimated cost by 58.7% in aggregate. Issue-type routing admits only a two-read prefix, whereas the two deeper workflows approach 75% fewer requests.

Table 1: Primary live-provider results ( $n = 30$  held-out records per family). “Interfaces” counts provider-visible tool interfaces. All conditions see identical records, model, and grader; *Manual* is hand-written with full knowledge of the workflow.

| Family                      | Exact held-out contract |              |        | Reduction of compiled vs. baseline (%) |            |        |         |      |
|-----------------------------|-------------------------|--------------|--------|----------------------------------------|------------|--------|---------|------|
|                             | Base                    | Compiled     | Manual | Requests                               | Interfaces | Tokens | Latency | Cost |
| Issue-type routing          | 30/30                   | <b>30/30</b> | 30/30  | 50.0                                   | 0.0        | 39.5   | 51.7    | 32.0 |
| PR-outcome audit            | 30/30                   | <b>30/30</b> | 30/30  | 75.0                                   | 66.7       | 80.8   | 73.0    | 75.3 |
| Backlog-attention routing   | 29/30                   | <b>30/30</b> | 30/30  | 74.8                                   | 66.3       | 81.4   | 68.9    | 75.1 |
| Weighted total ( $n = 90$ ) | 89/90                   | <b>90/90</b> | 90/90  | 66.6                                   | 44.2       | 63.1   | 64.2    | 58.7 |

Table 2: Where refusal happens, and what happens where it does not. Stage numbers are Alg. 1’s; “slots” is the share of dependency slots with a recovered producer, “p/a/w” means held-out replay pass/abstain/wrong, and  $n_{\max}$  is the largest family support against the 92 eq. (7) requires.

| Substrate | Traces | Slots (1) | Windows (2) | At floor (2) | Synth. (4) | Replay p/a/w (5)   | $n_{\max}$ (6) | Adm.     |
|-----------|--------|-----------|-------------|--------------|------------|--------------------|----------------|----------|
| NESTFUL   | 1,415  | 96.3%     | 1,207       | 32           | 12         | 24 / 12 / <b>0</b> | 26             | <b>0</b> |
| API-Bank  | 212    | 63.5%     | 37          | 8            | 2          | 0 / 2 / <b>0</b>   | 8              | <b>0</b> |
| BFCL v4   | 200    | 68.0%     | 86          | 9            | 4          | 3 / 3 / <b>0</b>   | 15             | <b>0</b> |
| AppWorld  | 136    | 63.2%     | 136         | 33           | 1          | 34 / 0 / <b>0</b>  | <b>136</b>     | <b>1</b> |

No held-out record is a compiled-only failure; the one-sided 95% upper bound on that discordance is 3.3%, which supports preservation on the observed cohort, not equivalence in general. Each artifact was admitted with 92 zero-violation calibration groups ( $U_{\eta} = 0.0498 \leq .05$ ). Hand-written programs also achieve 90/90 and are cheaper for issue-type routing.

## 5.2 TIME-FORWARD TRANSFER: COMPILE OR RETIRE

On a narrower two-read pull-request outcome task across frozen public repositories, with held-out records strictly newer than discovery records and selection sealed before any provider call, GAC completes 120/120 exact held-out pairs on four of five repositories and retires `pytorch/pytorch` at compile time rather than loosening its gate; requests fall 44.4% and tokens 52.4%. A balanced three-repository rerun reaches 180/180 exact pairs at 66.7% fewer requests and 78.4% fewer tokens (table 4). A fixed two-read template is more efficient in the five-repository core and essentially tied in the rerun, so this is again evidence of automatic discovery and principled retirement rather than of runtime dominance.

## 5.3 WHAT BINDS: SAMPLE SIZE, THEN DEPLOYABILITY

The external substrates isolate the binding constraint (table 2). All four clear provenance, synthesis, and held-out replay with *no* wrong execution anywhere. NESTFUL, API-Bank, and executed BFCL gold plans then retire, and each retirement was arithmetically settled before the compiler ran—which is why a fourth substrate was needed. AppWorld’s public gold solutions reach  $n_{\max} = 136$ , so it is the first external corpus where eq. (7) is satisfiable, and it admits one artifact at 92/92 zero-violation groups ( $U_{\eta} = 0.0498$ ) with 11/11 exact held-out replay, under both pre-registered effect arms. Two features bound it: the gate is step-like, and the admitted program is argument-free, so provenance fixed its *length* rather than its content—four further candidates retired, two on 89 and 90 eligible groups at  $U = 0.051$ , and one mined three calls but synthesized two because its login password slot is not groundable in every supporting group (appendix F.3).

Admissibility is not dispatchability. Gold solutions are not agent trajectories, so we also test the admitted region against AppWorld’s 28 released baseline runs. Over 8,190 held-out trajectories it is structurally eligible at the entry boundary in 37.9%, and that average is bimodal: near-total for full-code agents, and 1 of 4,680 for ReAct and plan-and-execute, which open by reading API documentation so the prefix leaves position 0 (table 5). These trajectories repeat tasks across 28 released runs and are descriptive units, not 8,190 independent samples. A miss abstains cleanly, and this

measurement bounds  $\phi$  from above rather than estimating a population dispatch rate, so no saving or inferential interval is claimed.

## 6 RELATED WORK

**Guarded trace and agent workflow compilation.** GAC adapts the tracing-JIT pattern—compile a hot path behind a guard and return to the interpreter on a miss (Bala et al., 2000; Hölzle et al., 1992)—to agents, adding grounded arguments and declared effects (Lucassen & Gifford, 1988). Agent systems mine meta-tools or workflows, cache plans, compile orchestration, validate plans, search graphs, or prune topology (Abuzakuk et al., 2026; Wang et al., 2025; Zhang et al., 2025; Wei et al., 2026; Winston et al., 2026; Li et al., 2026; Chen et al., 2026). DSPy optimizes prompts, GEPA evolves them, RouteLLM changes models, and LLMCompiler parallelizes current calls (Khatab et al., 2024; Agrawal et al., 2026; Ong et al., 2025; Kim et al., 2024). Unlike these targets, GAC treats recurrence only as a proposal: provenance, effects, position, contracts, and finite-sample risk each retain a veto.

**Synthesis, risk control, and evaluation.** Bounded search draws on programming by example, sketching, and provenance (Gulwani, 2011; Solar-Lezama et al., 2006; Cheney et al., 2009); admission combines Clopper–Pearson with fixed-grid learn-then-test (Clopper & Pearson, 1934; Angelopoulos et al., 2022). BFCL, ToolLLM, and AppWorld score tool use or interactive agents broadly (Patil et al., 2025; Qin et al., 2023; Trivedi et al., 2024). NESTFUL and API-Bank retain observed intermediate results directly; BFCL and AppWorld become post-trace compiler substrates only after official gold artifacts are executed on their pinned backends. None of those provider-free rows measures agent quality.

## 7 DISCUSSION AND LIMITATIONS

The supported claim is narrow: repeated read-only prefixes can become guarded programs that preserve held-out contracts on the tested GitHub families and retire when evidence is inadequate.

**Gate maturity is the main scientific risk.** Proposition 1 is per-candidate, not compiler-wide over Alg. 1’s adaptive search unless that search is frozen (corollary 2); the primary GitHub families do not freeze, so a multiplicity repair would raise their requirement further (appendix C). The gate is also empirically step-like at  $n_\eta \geq 92$ , AppWorld included: it *refuses* correctly far better than  $q$  *discriminates*, and refitting  $q$  for a genuine frontier is the key follow-up, for which a pre-registered but not-yet-run protocol already exists (appendix H).

**Group independence is assumed, not established.** Proposition 1 needs i.i.d. groups. Here each group is one record from one snapshot with `min_days=min_principals=1`, and AppWorld’s task-level groups share a scenario’s supervisor and database, so neither corpus establishes exchangeability and the second inherits the conditional rather than repairing it (appendix G). The 8,190 released AppWorld trajectories also repeat tasks across models, architectures, and runs; their structural-eligibility rate is a descriptive deployment diagnostic, not an independent-sample estimate.

**External validity, dispatchability, and competitiveness.** The richer evidence comes from one revision-pinned repository and one provider/model family, and the cross-repository extension is a narrower, template-tied task. Admission is also not dispatchability: a prefix is only worth compiling for an architecture that issues one, which neither proposition 1 nor the guard tests. We establish no superiority under drift or across providers and do not price manual authoring or maintenance, though a live compression-comparator ablation and a pre-registered drift-parity protocol narrow both gaps (appendix G). Finally, exact task contracts do not certify a downstream model continuation, and writes, handoffs, hosted tools, undeclared effects, and removed-turn guardrails lie outside proposition 1 entirely (appendix G).

## 8 CONCLUSION

GAC treats recurrence as candidate evidence, not permission to rewrite an agent: it preserves exact held-out contracts while saving resources, retires all families on three external substrates, and on the

fourth admits a small region, retires four more at the margin, and is structurally eligible only under some agent architectures. Judge workflow optimizers by their evidence, refusals, and fallback, not only by the turns they remove.

## AI USE STATEMENT

The authors used generative AI tools to restructure and copy-edit the manuscript, suggest section organization, draft and revise paper text, critique methodological exposition and internal consistency, and convert repository-grounded results into LaTeX tables. They were not used to generate experimental results or to replace manual verification. Every reported number, citation, equation, and comparison was checked against the repository artifacts and retained outputs, and the authors take responsibility for the final content.

## ETHICS STATEMENT

The main risk in removing model-mediated decisions is over-trusting a compiled prefix, and thereby accelerating unsafe automation or deleting a guardrail evaluation that happened to run at the removed boundary. The method mitigates this by compiling only read-only prefixes, treating unknown effects, approvals, handoffs, and writes as hard barriers, and defaulting to retirement or fallback when evidence is insufficient. All study data are public records (GitHub repository metadata) or public benchmarks; no personal or private data were collected. The artifact is not a license to bypass human review, legal or compliance controls, or safety guardrails; unsupported regimes are part of the result, not exceptions to hide.

## REPRODUCIBILITY STATEMENT

Algorithms 1 and 2 give the compiler and admission protocols, Algs. 3 and 4 give provenance construction and runtime dispatch, appendix A proves proposition 1, and appendix C records the registered statistical and synthesis settings. The accompanying artifact contains source manifests, raw result files, deterministic table and figure generation, and scripts for the live GitHub studies, the cross-repository extension, and all four compiler substrates: NESTFUL, API-Bank, executed BFCL, and executed AppWorld. Appendix F.3 states how the latter two obtain observed results, and appendix C gives their exact provider-free entry points. Before upload, these materials will be supplied as an anonymous archive with hashes and manifests preserved.

## REFERENCES

- Sami Abuzakuk, Anne-Marie Kermarrec, Rishi Sharma, Rasmus Moorits Veski, and Martijn de Vos. Optimizing agentic workflows using meta-tools. *arXiv preprint arXiv:2601.22037*, 2026. URL <https://arxiv.org/abs/2601.22037>.
- Lakshya A. Agrawal, Shangyin Tan, Dilara Soylu, Noah Ziems, Rishi Khare, Krista Opsahl-Ong, Arnav Singhvi, Herumb Shandilya, Michael J. Ryan, Meng Jiang, Christopher Potts, Koushik Sen, Alexandros G. Dimakis, Ion Stoica, Dan Klein, Matei Zaharia, and Omar Khattab. GEPA: Reflective prompt evolution can outperform reinforcement learning. In *International Conference on Learning Representations*, 2026. URL <https://arxiv.org/abs/2507.19457>. Oral presentation.
- Anastasios N. Angelopoulos, Stephen Bates, Emmanuel J. Candès, Michael I. Jordan, and Lihua Lei. Learn then test: Calibrating predictive algorithms to achieve risk control. In *International Conference on Learning Representations*, 2022. URL <https://openreview.net/forum?id=Tnc0rox3tQ>.
- Vasanth Bala, Evelyn Duesterwald, and Sanjeev Banerjia. Dynamo: A transparent dynamic optimization system. In *Proceedings of the ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI)*, pp. 1–12, 2000. doi: 10.1145/349299.349303.

- Kinjal Basu, Ibrahim Abdelaziz, Kiran Kate, Mayank Agarwal, Maxwell Crouse, Yara Rizk, Kelsey Bradford, Asim Munawar, Sadhana Kumaravel, Saurabh Goyal, Xin Wang, Luis A. Lastras, and Pavan Kapanipathi. NESTFUL: A benchmark for evaluating LLMs on nested sequences of API calls. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pp. 33538–33547, 2025. doi: 10.18653/v1/2025.emnlp-main.1702. URL <https://aclanthology.org/2025.emnlp-main.1702/>.
- Yulang Chen, Haoxuan Peng, Jinyan Liu, Zichen Wen, Dongrui Liu, and Linfeng Zhang. AgentSlimming: Towards efficient and cost-aware multi-agent systems. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics*, pp. 30064–30086, 2026. doi: 10.18653/v1/2026.acl-long.1387. URL <https://aclanthology.org/2026.acl-long.1387/>.
- James Cheney, Laura Chiticariu, and Wang-Chiew Tan. Provenance in databases: Why, how, and where. *Foundations and Trends in Databases*, 1(4):379–474, 2009. doi: 10.1561/19000000006.
- C. J. Clopper and E. S. Pearson. The use of confidence or fiducial limits illustrated in the case of the binomial. *Biometrika*, 26(4):404–413, 1934. doi: 10.1093/biomet/26.4.404.
- Sumit Gulwani. Automating string processing in spreadsheets using input-output examples. In *Proceedings of the 38th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, pp. 317–330, 2011. doi: 10.1145/1926385.1926423.
- Headroom Labs. Headroom: Compressing tool outputs, logs, and retrieved context before they reach the LLM. Software repository, <https://github.com/headroomlabs-ai/headroom>, 2026. Apache-2.0; self-reported 60–95% token reduction on JSON payloads and 73% on a GitHub triage workload. Accessed 7 August 2026.
- Urs Hölzle, Craig Chambers, and David Ungar. Debugging optimized code with dynamic deoptimization. In *Proceedings of the ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI)*, pp. 32–43, 1992. doi: 10.1145/143095.143114.
- Omar Khattab, Arnav Singhvi, Paridhi Maheshwari, Zhiyuan Zhang, Keshav Santhanam, Sri Vardhamanan, Saiful Haq, Ashutosh Sharma, Thomas T. Joshi, Hanna Moazam, Heather Miller, Matei Zaharia, and Christopher Potts. DSPy: Compiling declarative language model calls into self-improving pipelines. In *International Conference on Learning Representations*, 2024. URL <https://openreview.net/forum?id=sY5N0zY50d>.
- Sehoon Kim, Suhong Moon, Ryan Tabrizi, Nicholas Lee, Michael W. Mahoney, Kurt Keutzer, and Amir Gholami. An LLM compiler for parallel function calling. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pp. 24370–24391, 2024. URL <https://proceedings.mlr.press/v235/kim24y.html>.
- Junyan Li, Zhang-Wei Hong, Maohao Shen, Yang Zhang, and Chuang Gan. FlowCompile: An optimizing compiler for structured LLM workflows. *arXiv preprint arXiv:2605.13647*, 2026. URL <https://arxiv.org/abs/2605.13647>.
- Minghao Li, Yingxiu Zhao, Bowen Yu, Feifan Song, Hangyu Li, Haiyang Yu, Zhoujun Li, Fei Huang, and Yongbin Li. API-bank: A comprehensive benchmark for tool-augmented LLMs. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pp. 3102–3116, 2023. doi: 10.18653/v1/2023.emnlp-main.187. URL <https://aclanthology.org/2023.emnlp-main.187/>.
- John M. Lucassen and David K. Gifford. Polymorphic effect systems. In *Proceedings of the 15th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL)*, pp. 47–57, 1988. doi: 10.1145/73560.73564.
- Isaac Ong, Amjad Almahairi, Vincent Wu, Wei-Lin Chiang, Tianhao Wu, Joseph E. Gonzalez, M. Waleed Kadous, and Ion Stoica. RouteLLM: Learning to route LLMs with preference data. In *International Conference on Learning Representations*, 2025. URL [https://proceedings.iclr.cc/paper\\_files/paper/2025/hash/5503a7c69d48a2f86fc00b3dc09de686-Abstract-Conference.html](https://proceedings.iclr.cc/paper_files/paper/2025/hash/5503a7c69d48a2f86fc00b3dc09de686-Abstract-Conference.html).

- Shishir G. Patil, Huanzhi Mao, Fanjia Yan, Charlie Cheng-Jie Ji, Vishnu Suresh, Ion Stoica, and Joseph E. Gonzalez. The berkeley function calling leaderboard (BFCL): From tool use to agentic evaluation of large language models. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pp. 48371–48392, 2025. URL <https://proceedings.mlr.press/v267/patil25a.html>.
- Yujia Qin, Shihao Liang, Yining Ye, Kunlun Zhu, Lan Yan, Yaxi Lu, Yankai Lin, Xin Cong, Xiangru Tang, Bill Qian, Sihan Zhao, Lauren Hong, Runchu Tian, Ruobing Xie, Jie Zhou, Mark Gerstein, Dahai Li, Zhiyuan Liu, and Maosong Sun. ToolLLM: Facilitating large language models to master 16000+ real-world APIs. *arXiv preprint arXiv:2307.16789*, 2023. URL <https://arxiv.org/abs/2307.16789>.
- Armando Solar-Lezama, Liviu Tancau, Rastislav Bodik, Sanjit Seshia, and Vijay Saraswat. Combinatorial sketching for finite programs. In *Proceedings of the 12th International Conference on Architectural Support for Programming Languages and Operating Systems*, pp. 404–415, 2006. doi: 10.1145/1168857.1168907.
- Harsh Trivedi, Tushar Khot, Mareike Hartmann, Ruskin Manku, Vinty Dong, Edward Li, Shashank Gupta, Ashish Sabharwal, and Niranjan Balasubramanian. AppWorld: A controllable world of apps and people for benchmarking interactive coding agents. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 16022–16076, 2024. doi: 10.18653/v1/2024.acl-long.850. URL <https://aclanthology.org/2024.acl-long.850/>.
- Zora Zhiruo Wang, Jiayuan Mao, Daniel Fried, and Graham Neubig. Agent workflow memory. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pp. 63897–63911, 2025. URL <https://proceedings.mlr.press/v267/wang25bx.html>.
- Lei Wei, Qi Liu, Ruiyang Huang, Xiao Peng, Sicong Xie, Lanbo Lin, Chenhao Jiang, Yuanwu Xu, Tianyuan Yang, Jiayao Liu, Li Cai, Zhaolu Kang, and Bin Wang. EvoC2F: Compiling tool orchestration for efficient and evolvable LLM agents. In *Proceedings of the 43rd International Conference on Machine Learning*, volume 306 of *Proceedings of Machine Learning Research*, 2026. URL <https://openreview.net/forum?id=ZSGB91kMOG>.
- Caleb Winston, Ron Yifeng Wang, Azalia Mirhoseini, and Christos Kozyrakis. Agent JIT compilation for latency-optimizing web agent planning and scheduling. In *Proceedings of the 43rd International Conference on Machine Learning*, volume 306 of *Proceedings of Machine Learning Research*, 2026. URL <https://arxiv.org/abs/2605.21470>.
- Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. ReAct: Synergizing reasoning and acting in language models. In *International Conference on Learning Representations*, 2023. URL [https://openreview.net/forum?id=WE\\_vluYUL-X](https://openreview.net/forum?id=WE_vluYUL-X).
- Qizheng Zhang, Michael Wornow, and Kunle Olukotun. Cost-efficient serving of LLM agents via test-time plan caching. *arXiv preprint arXiv:2506.14852*, 2025. URL <https://arxiv.org/abs/2506.14852>.

## A PROOF OF PROPOSITION 1

Condition on the train/development data, fix  $\eta \in \Lambda$ , and write  $\gamma = \delta/|\Lambda|$ . Let  $D_i^\eta$  and  $W_i^\eta$  be the group admission and violation indicators defined in section 2. Because the entire candidate and grid were frozen before calibration and groups are i.i.d., the pairs  $(D_i^\eta, W_i^\eta)$  are i.i.d. Conditional on  $n_\eta = m > 0$  admitted groups, their violation count therefore satisfies  $k_\eta \mid n_\eta = m \sim \text{Binomial}(m, r_\eta)$ , where  $r_\eta = \Pr(W^\eta = 1 \mid D^\eta = 1)$ . Inverting the one-sided exact binomial test (Clopper & Pearson, 1934) gives the upper bound  $U_\eta = \text{Beta}^{-1}(1 - \gamma; k_\eta + 1, m - k_\eta)$  of eq. (7), which satisfies  $\Pr\{r_\eta > U_\eta \mid n_\eta = m\} \leq \gamma$ . When  $m = 0$  the event is impossible under the convention that empty-admission thresholds receive the vacuous bound  $U_\eta = 1$ . Averaging over the random value of  $n_\eta$  therefore yields  $\Pr\{r_\eta > U_\eta\} \leq \gamma$  for each fixed threshold. Applying the union bound over the finite grid gives  $\Pr\{\exists \eta \in \Lambda : r_\eta > U_\eta\} \leq |\Lambda|\gamma = \delta$ , i.e. simultaneous coverage with probability at least  $1 - \delta$ . On that event, every threshold whose upper bound is at most  $\alpha$  satisfies the group-level target in eq. (4), including the threshold returned by Alg. 2. ■

Two scope conditions are worth restating. First, the statement is *per fixed candidate*: Alg. 1 may calibrate several families and retain survivors, and that outer search is not multiplicity-corrected — except when the search is frozen, per the corollary below. Second, i.i.d. source groups is an assumption about the calibration cohort, not a property the compiler verifies; the grouped split in Alg. 1 makes it plausible by construction but does not establish it under temporal drift.

**Corollary 2** (Compiler-wide guarantee under frozen selection). *If Alg. 1 runs with candidate freezing enabled — ranking on train data, synthesizing and challenging on train and development windows only, and halting the search the instant any candidate reaches calibration, whatever its outcome — proposition 1’s guarantee holds for the run’s actual reported output with no adjustment to  $\gamma = \delta/|\Lambda|$ : no candidate multiplicity remains to correct for.*

*Proof.* Ranking, synthesis, the challenge test, and the frozen score model  $q$  are all functions of train and development windows alone, fixed before a single calibration window is inspected. The search reaches calibration for the top-ranked surviving candidate exactly once and halts immediately after that candidate’s outcome, so no second candidate’s program, guard, verifier, or score model is ever compared against calibration data. Proposition 1’s precondition — one candidate fixed before observing calibration — is met by the run’s entire search, not by a hypothetical single-candidate sub-procedure, so its conclusion transfers unchanged. ■ □

This is a corollary, not a new inequality: it is a restatement of proposition 1 for a search that never lets a second candidate see calibration data. It is not hypothetical: the cross-repository extension already runs with freezing enabled, and every sealed repository report there shows exactly one candidate reaching calibration, so corollary 2 applies to those admissions directly. The three primary GitHub families do not use it — each mines two candidates that both reach calibration, with the higher-removal survivor kept by dominance — so they remain the per-candidate certificates above and are not retroactively strengthened. `paper/scripts/frozen_candidate_coverage_simulation.py` backs this with an exact (non-Monte-Carlo) combinatorial check: at  $n = 92$ , the population rate maximizing a single candidate’s miscoverage is  $r^* \approx 0.637$  at miscoverage  $\approx 0.00909 \approx \gamma$ ; pooling  $m$  such candidates at the same per-candidate budget with no correction lets the probability that *some* certificate is wrong grow with  $m$  (1.8% at  $m = 2$  to 13.6% at  $m = 16$ , crossing  $\delta = .10$  past the eleven-point grid), while the Bonferroni correction already discussed keeps every tested  $m$  under 1%, at the group cost this appendix already reports (106 at  $m = 2$ , recomputed here independently). A companion seeded Monte Carlo check (seed 20260822,  $M=200,000$ ) draws the 92 groups in correlated blocks instead of i.i.d. and finds realized miscoverage matches the exact baseline at zero correlation but reaches 56% against a 0.9% nominal budget at the most clustered configuration tested, quantifying rather than only asserting why the live studies’ `min_days=min_principals=1` configuration is this guarantee’s weakest link.

## B ADDITIONAL ALGORITHMS

Algorithm 3 constructs the typed provenance graph that decides eq. (1). Its two block conditions are asymmetric on purpose: a slot with no witness is treated as a genuine model decision and a slot with

**Algorithm 3** BUILDPATG — typed argument provenance. A slot with no witness is a genuine model decision; a slot with too many witnesses is ambiguous. Both block the region rather than defaulting to the cheapest explanation.

---

**Input:** episodes  $\mathcal{T}$ ; entry schema  $\Theta$ ; effect catalog  $C$ ; transform library  $\mathcal{L}$ ; ambiguity cap  $\kappa$   
**Output:** provenance graph  $G$  retaining  $1, \dots, \kappa$  candidate edges per nonliteral grounded slot; declared literals are marked separately

---

```

1:  $G \leftarrow \emptyset$ 
2: for all episodes  $E = (z, M, e_{1:T}, y) \in \mathcal{T}$  do
3:    $\Sigma \leftarrow \{("z", \pi, v) : (\pi, v) \in \text{FLATTEN}(z), \pi \in \Theta\}$  ▷ only allowlisted entry paths are eligible
4:   for all tool calls  $c_j \in e_{1:T}$  in order do
5:     for all argument slots  $u \in \text{args}(c_j)$  do
6:       if  $u$  is declared literal-only for  $c_j$  then
7:         mark  $(c_j, u)$  LITERAL; continue ▷ preserve the declared constant; infer no provenance
8:       end if
9:        $\mathcal{H} \leftarrow \text{BESTPERSOURCE}\{(\sigma, g) \in \Sigma \times \mathcal{L} : g(\text{value}(\sigma)) = u, \text{depth}(g) \leq 2\}$ 
10:      if  $\mathcal{H} = \emptyset$  then
11:        mark  $(c_j, u)$  MODEL-ORIGINATED; block region
12:      else if  $|\mathcal{H}| > \kappa$  then
13:        mark  $(c_j, u)$  AMBIGUOUS; block region
14:      else
15:         $G \leftarrow G \cup \{\sigma \xrightarrow{g} (c_j, u) : (\sigma, g) \in \mathcal{H}\}$  ▷ synthesis chooses a stable source across groups
16:      end if
17:    end for
18:     $\Sigma \leftarrow \Sigma \cup \{(\text{var}(c_j), \pi, v) : (\pi, v) \in \text{FLATTEN}(o_j), \text{ELIGIBLE}(\pi, v)\}$  ▷ a result witnesses only later slots
19:  end for
20:  for all pairs  $(e_s, e_t)$  sharing a resource where either writes it do
21:     $G \leftarrow G \cup \{e_s < e_t\}$  ▷ order edge: no reordering is ever proposed across a conflict
22:  end for
23: end for
24: return  $G$ 

```

---

more than  $\kappa$  witnesses is treated as ambiguous, so the analysis never resolves an argument by picking the cheapest available explanation.

Algorithm 4 is the runtime counterpart of eq. (3). Cheap deterministic checks reject first and the calibrated gate runs only on what survives them, so the common case costs a registry lookup and a guard evaluation. Exactly one exit path compacts; every clean rejection returns the unmodified baseline agent, and any failure whose reversibility cannot be attested raises an incident rather than feigning a rollback. “Baseline” is byte-exact only where a staging owner holds the commit boundary: the outer runner can restore byte-identical model input, whereas a model-boundary adapter cannot un-emit a response the host has already committed to conversation history.

## C CONFIGURATION

All reported admission runs use risk budget  $\alpha = 0.05$ , confidence budget  $\delta = 0.1$ , and a fixed grid of  $|\Lambda| = 11$  thresholds, which fixes the zero-violation sample requirement at  $n_\eta \geq \log(\delta/|\Lambda|)/\log(1 - \alpha) = 91.6$ , i.e. 92 admitted groups; at exactly  $n_\eta = 92$  and  $k_\eta = 0$  the bound is  $U_\eta = 1 - (0.1/11)^{1/92} = 0.0498 \leq \alpha$ . Tightening the confidence budget to  $\delta = 0.05$  would raise the requirement to 106 groups, which no study in this paper reaches, so  $\delta = 0.1$  is the operative setting throughout and is reported as such rather than rounded to match  $\alpha$ . A Bonferroni repair over  $m = 2$  candidate families would use  $\gamma = \delta/(m|\Lambda|)$  and land on the same 106-group requirement, which is why the paper reports per-candidate certificates rather than a compiler-wide one. The grid is  $\Lambda = \{.02, .05, .08, .11, .14, .17, .20, .25, .30, .40, .50\}$  and the ambiguity cap is  $\kappa = 3$ . Synthesis uses a 23-operator closed library at depth  $\leq 2$ ; region mining uses  $w_{\max}$  tool calls per window and requires support across independent scenario groups. The implementation also supports a minimum-distinct-days requirement, which the single-snapshot live study does not exercise (`min_days` = 1). Two terms of eq. (5) are coarser in the implementation than the notation suggests:  $\rho_{\text{eff}}$  is approximated by the fraction of tool names containing a namespace separator rather than read from the effect catalog, and

---

**Algorithm 4** DISPATCH — boundary-time admission. Cheap deterministic checks reject first and the calibrated gate runs only on what survives them. Exactly one exit path compacts; clean rejections return the unmodified agent, while an unattested or post-commit failure raises an incident.

---

**Input:** entry state  $z$ ; runtime context  $M'$ ; registry  $\mathcal{R}$ ; effect catalog  $C$ ; mode  $m \in \{\text{off}, \text{shadow}, \text{live}\}$

**Output:** observations for the region, FALLBACK to baseline agent  $B$ , or INCIDENT

---

```

1: if  $m = \text{off}$  then
2:   return FALLBACK                                ▷ conformance: model-visible input must be byte-identical
3: end if
4:  $A \leftarrow \mathcal{R}.\text{RESOLVE}(M'.\text{compatibility}, M'.\text{partition})$ 
5: if  $A = \perp \vee A.\text{lifecycle} \neq \text{ACTIVE} \vee \neg \text{VERIFYSIGNATURE}(A)$  then
6:   return FALLBACK
7: end if
8: if  $H_A(z, M') \neq 1$  then
9:   return FALLBACK                                ▷ manifest pins, isolation keys, typed hulls, effect set
10: end if
11: if  $\text{tools}(A.P) \cap \text{ALREADYOBSERVED} \neq \emptyset$  then
12:   return FALLBACK                                ▷ the region already ran; its live-ins are not the fitted ones
13: end if
14: if  $m = \text{live} \wedge \exists f \in \text{tools}(A.P) : C(f).\text{quota\_attested} \wedge \text{SNAPSHOT} = \perp$  then
15:   return FALLBACK                                ▷ no staging owner, so a discarded attempt is not attestable
16: end if
17: if  $q_A(z) > \eta_A$  then
18:   return FALLBACK                                ▷ calibrated abstention on an uncertain input
19: end if
20: if  $m = \text{shadow}$  then
21:   log would-dispatch; return FALLBACK            ▷ observe coverage without changing behaviour
22: end if
23:  $S \leftarrow \text{STAGE.BEGIN}(); r \leftarrow \text{INTERPRET}(A.P, z, \text{FACADE}(C))$   ▷ the facade re-checks effects at execution
   time
24: if  $r$  raised PreCommitError then
25:   return FALLBACK if  $S.\text{ABORTANDATTEST}()$ , else INCIDENT
26: else if  $r$  raised PostCommitError then
27:   return INCIDENT                                ▷ an external commitment happened; do not feign rollback
28: end if
29: if  $V_A(r) \neq 1$  then
30:   return FALLBACK if  $S.\text{ABORTANDATTEST}()$ , else INCIDENT  ▷ live-outs, effect multiset, and call count
   are checked before use
31: end if
32: return  $r.\text{observations}$  if  $S.\text{COMMIT}(r.\text{effects})$  succeeds, else INCIDENT

```

---

$|F|$  is substituted by the argument-slot count of the first observed window because ranking precedes synthesis. Both appear only in a ranking heuristic — no admission decision depends on either, and the effect catalog is consulted where it is load-bearing, in qualification and in the runtime facade.

The two executable-gold external substrates are provider-free after their pinned sources are prepared. Their exact compiler entry points are:

```

B=paper/scripts/bfcl_compiler_benchmark.py
python "$B" --source-root "$BENCHMARK_SOURCE_ROOT"
A=paper/scripts/appworld_compiler_benchmark.py
python "$A" --source-root "$BENCHMARK_SOURCE_ROOT" \
  --appworld-python "$APPWORLD_VENV/bin/python"
D=paper/scripts/appworld_dispatch_preflight.py
python "$D" --source-root "$BENCHMARK_SOURCE_ROOT"

```

The AppWorld environment uses Python 3.12 with package version 0.1.3.post1 and data version 0.1.0; the signed source manifest and requirements file pin the remaining acquisition inputs. The BFCL source is pinned to revision 6ea57973c7a6. These commands regenerate only aggregate, non-sensitive evidence; no model or provider call is made.

Table 3: Primary GitHub benchmark at a glance. Each sample supplies one record number from a pinned public snapshot; every listed answer field is checked against the source record.

| Task              | Read-only evidence path                                              | Exact answer contract                                                                |
|-------------------|----------------------------------------------------------------------|--------------------------------------------------------------------------------------|
| Issue type        | <code>issue_number</code> → record, official labels, comments        | Title, state, label-derived category, verbatim excerpt                               |
| PR outcome        | <code>pr_number</code> → PR record, merge status, discussion         | open, merged, or closed_unmerged, plus title and comment excerpt                     |
| Backlog attention | <code>issue_number</code> → open-issue record, assignees, discussion | owned, discussed_unowned, or awaiting_first_response, plus title, owner, and excerpt |

## D PAIRED STATISTICS FOR THE PRIMARY FAMILIES

Per-record paired differences (compiled minus baseline) on the two families that admit a deep prefix, with 10,000-sample bootstrap intervals and two-sided Wilcoxon signed-rank tests over  $n = 30$  paired held-out records:

| Family     | Endpoint       | Paired difference (95% CI) | $p$                  |
|------------|----------------|----------------------------|----------------------|
| PR-outcome | Total tokens   | -2196.7 [-2211.1, -2182.5] | $1.7 \times 10^{-6}$ |
|            | Wall latency   | -3889 ms [-4499, -3346]    | $1.9 \times 10^{-9}$ |
|            | Estimated cost | $-\$5.11 \times 10^{-4}$   | $1.7 \times 10^{-6}$ |
| Backlog    | Total tokens   | -2315.0 [-2366.0, -2224.1] | $1.7 \times 10^{-6}$ |
|            | Wall latency   | -4428 ms [-5400, -3558]    | $1.9 \times 10^{-9}$ |
|            | Estimated cost | $-\$5.44 \times 10^{-4}$   | $1.7 \times 10^{-6}$ |

Provider-request differences are deterministic by construction (every pair removes the same number of turns), so no rank test is reported for them: the statistic would measure the degeneracy rather than the effect.

## E ADDITIONAL EXPERIMENTAL DETAIL

**Primary workflow families.** The three-family result is the main paper’s strongest live evidence because it combines distinct tool vocabularies, exact task graders, balanced held-out cohorts, and provider-backed execution. The issue-type family intentionally retains a shallower two-read prefix; the newer pull-request and backlog families admit deeper pre-model compaction. This difference explains the lower resource savings on the first family and is part of the evidence boundary, not noise to average away.

**Cross-repository extension.** The cross-repository study is narrower than the primary workflow families. Its task uses two read-only tools and a simpler exact outcome contract, which is precisely why a fixed template can remain tied or even more efficient than the learned artifact. We keep the result in the main text because it addresses a reviewer-critical scope question—whether guarded specialization survives beyond one repository—while being explicit about the reduced task complexity.

**The position invariant, empirically.** An early pilot dispatched a compiled *suffix* region at the entry boundary, violating eq. (2). Because the runtime resolves artifacts before the agent has made the calls the region assumed had already happened, the pilot reordered and duplicated tool calls. The failure is a semantic mismatch between compilation-time position and resolution-time position, not a tuning error, which is why eq. (2) is a hard constraint in the deployed compiler rather than a scoring penalty.

**Omitted from the main claim.** The repository retains additional studies that are useful scientifically but less suitable for the main narrative:

- a natural-order issue study that exposed a factual continuation miss for a more aggressive three-read artifact,

Table 4: Disaggregated cross-repository time-forward outcomes behind section 5.2. The core protocol uses four completed repositories and one compile-time retirement; the balanced rerun changes the selected repositories and class mix, so it is reported separately rather than pooled with the core protocol. Reductions compare the compiled condition with the paired unchanged baseline.

| Protocol | Repository           | Held-out | Exact  | Req. ↓ | Tokens ↓ |
|----------|----------------------|----------|--------|--------|----------|
| Core     | huggingface/datasets | 30       | 30/30  | 44.4%  | 52.2%    |
| Core     | pandas-dev/pandas    | 30       | 30/30  | 44.4%  | 52.5%    |
| Core     | psf/requests         | 30       | 30/30  | 44.4%  | 52.3%    |
| Core     | streamlit/streamlit  | 30       | 30/30  | 44.4%  | 52.7%    |
| Core     | pytorch/pytorch      | —        | RETIRE | —      | —        |
| Balanced | pandas-dev/pandas    | 60       | 60/60  | 66.7%  | 78.6%    |
| Balanced | psf/requests         | 60       | 60/60  | 66.7%  | 78.7%    |
| Balanced | pytorch/pytorch      | 60       | 60/60  | 66.7%  | 78.0%    |

- exploratory guarded composite synthesis and a follow-up comparison to a fair pre-model manual program,
- a bounded GEPA prompt-optimization study in which GEPA retained its seed, so the combined arm is a replication rather than evidence of interaction, and
- a portfolio pilot and controlled demonstration suite for runtime boundary coverage, including workloads where removing turns *increases* cost through prefix-cache fragmentation.

These studies remain important to the broader artifact, but they either use a weaker risk level, address a narrower engineering question, or are explicitly exploratory.

## F EVIDENCE BREADTH AND BOUNDARY AUDITS

The main text uses the nine-page budget for the registered live studies, the time-forward extension, and the four compiler-measured external substrates. This appendix makes the retained breadth visible without pooling incomparable cohorts or promoting screened and exploratory work to deployment evidence.

### F.1 REPOSITORY-LEVEL TIME-FORWARD OUTCOMES

The core result succeeds on four repositories (120/120 held-out pairs) and retains the fifth as a compile-time retirement after 116 exact discovery traces. The balanced rerun has a distinct three-repository cohort and reaches 180/180 held-out pairs. These rows show breadth across repositories, not independent evidence for an average effect: all runs share the same provider family and the same narrow two-read outcome task.

### F.2 CONTINUATION AUDIT

The earlier natural-order three-read study is excluded from the registered  $\alpha = .05$  headline because it used  $\alpha = .10$  and exposed a downstream continuation error. Its retained counterfactual replay nevertheless tests a boundary the exact tool contract cannot see: 17/18 candidate continuations pass unchanged, while issue 6602 has a comment-evidence mismatch. Checked rendering repairs that one case, yielding 18/18 final contracts. This audit makes no provider calls and does not estimate live agent quality; it establishes only that the continuation check caught a concrete failure the compacted tool trace alone would miss.

### F.3 EXTERNAL EVIDENCE MAP

Table 6 records every named external benchmark in the retained evidence register and its admissible interpretation.

**Obtaining observed results from BFCL.** BFCL’s reference plans carry no tool results, which is why the plan-format row and the compiler row are listed separately. Its multi-turn categories, how-

Table 5: Structural dispatch eligibility of the admitted AppWorld artifact on the 28 released official baseline runs. Eligibility is the admitted program’s exact call sequence at normalized position 0; it is an upper bound on  $\phi$ , because the manifest check and the calibrated gate are not evaluated and the released logs retain API calls rather than model boundaries. Trajectories repeat tasks across runs and are descriptive rather than independent samples. A miss abstains to the unchanged agent.

| Released baseline architecture      | Runs | Trajectories | Eligible at pos. 0 | Rate  |
|-------------------------------------|------|--------------|--------------------|-------|
| Full code + reflection              | 8    | 2,340        | 2,339              | 1.000 |
| Iterative parallel function calling | 4    | 1,170        | 762                | 0.651 |
| Plan and execute                    | 8    | 2,340        | 1                  | 0.000 |
| ReAct                               | 8    | 2,340        | 0                  | 0.000 |
| Pooled                              | 28   | 8,190        | 3,102              | 0.379 |

ever, ship an executable stateful backend: each task publishes an `initial_config` snapshot and the classes it involves, so executing the *official gold plan* through the upstream helper yields observed results, an entry snapshot, and an in-process replay oracle, with no model in the loop. Five preconditions fail the run closed rather than publishing: an undeclared method, any upstream execution error, a read-like declaration observed mutating state, a compilable declaration observed advancing the scenario RNG, and any inexact re-execution. All held: 200 tasks and 1,142 calls executed cleanly and a second independent pass reproduced every result byte for byte, executing turn-at-a-time where the first pass executed call-at-a-time.

**Why the catalog is signed rather than inferred.** The same run audits the declarations empirically by snapshotting every involved instance around every call. It observed 1,060 state-mutating calls over 44 of the 81 methods and 94 calls that advanced the seeded scenario RNG, with zero cases of a read-like declaration mutating state or a compilable declaration advancing the RNG. One case is worth naming: a method whose name and docstring present it as a cost getter mutated its instance’s internal lookup on all 36 of its calls and is therefore declared a write. A name- or docstring-derived effect label would have admitted a write into a compiled region. Two further methods are read-like but carry no capabilities, because one derives an age from the host date and the other draws from — and advances — the seeded RNG; they are barriers by declaration and account for 35 suppressed spans. Write barriers suppress 2,802 spans in total, against 45 for entry-schema admissibility, 31 for slot ambiguity, and 5 for partial runs.

**Obtaining observed results from AppWorld, and why it is the reachable case.** AppWorld’s shipped `api_calls.json` records method, url, and request data but no responses, so it fails closed for the same reason BFCL’s plans do. Executing the official `compiled_solution_code` against the pinned `0.1.3.post1` backend supplies the results; each task’s own `specs.json` supplies the entry snapshot; a second execution supplies the replay oracle. Calls are captured at the single dispatch point every `apis.<app>.<api>(...)` call passes through, so the recorded identity is a semantic app/API pair with typed keyword arguments and a decoded JSON result. Of the 147 public train+dev gold solutions, 136 execute cleanly; 11 dev solutions raise against the shipped data version and are reported as upstream compatibility outcomes rather than dropped. The second pass reproduced 7,070/7,070 results byte for byte, minted access tokens included, because AppWorld freezes each task’s clock.

The empirical audit attributes every `INSERT/UPDATE/DELETE/CREATE/DROP` statement to the enclosing call. It observed 1,320 mutating calls, zero read-like declarations mutating, and—unlike BFCL—zero APIs that mutate on some calls but not others. Every `GET` was clean on every call and every `DELETE` and `PATCH` mutated on every call, so the effect boundary is crisp; the five `*.login` endpoints are the only non-`GET` APIs that never mutate, which is why they are the single contestable declaration and why the run is pre-registered in two arms. Groups are tasks, not scenarios: AppWorld’s three variants of a scenario share a supervisor and a database, so the bound is conditional on task-level independence, which this corpus does not establish—the same conditional the live studies carry, neither weakened nor repaired here. The calibration share is fixed at exactly the 92 groups the bound needs, since a larger set would make the bound easier and a smaller one would make it unreachable.

Table 6: Disposition of every retained external benchmark family. Only the four compiler-measurement rows support compiler claims; the remaining rows document why broader benchmark names are not silently converted into GAC results. BFCL appears twice because its reference plans and its executed gold plans are different evidence.

| Benchmark           | Evidence tier        | Retained result / claim boundary                                                                                                                         |
|---------------------|----------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------|
| NESTFUL             | Compiler measurement | 1,415 traces; 24 / 12 / 0 replay; all families retire.                                                                                                   |
| API-Bank            | Compiler measurement | 212 traces; 0 / 2 / 0 replay; all candidates retire.                                                                                                     |
| BFCL v4 (plans)     | Gold-plan checker    | 200/200 valid plans; plan format only.                                                                                                                   |
| BFCL v4 (executed)  | Compiler measurement | 200 traces, 1,142/1,142 exact re-execution; 3 / 3 / 0 replay; all families retire.                                                                       |
| AppWorld (executed) | Compiler measurement | 136/147 traces, 7,070/7,070 exact re-execution; $n_{\max} = 136 \geq 92$ ; one artifact admitted at $U = 0.0498$ ; four candidates retire at the margin. |
| ToolSandbox         | Live-provider subset | One simulated scenario; compiler not run.                                                                                                                |
| $\tau^2/\tau^3$     | Live-provider subset | 0/4 pass; compiler not run.                                                                                                                              |
| ToolBench           | Adapter smoke        | Ten fixtures; full data/backend unavailable.                                                                                                             |
| AgentBench          | Adapter/preflight    | Services and Freebase data gated.                                                                                                                        |
| GAIA                | Access preflight     | Authorization denied; no metric.                                                                                                                         |
| SWE-bench Verified  | Task adapter         | Official run host-gated.                                                                                                                                 |
| BrowseComp          | Live-web subset      | 1/3 correct, 28 searches; bypassed by the compiler.                                                                                                      |

Table 7: Gate-behaviour inventory from the retained threshold analysis. A partial frontier would show graded coverage as the threshold changes; none appears among the seven current registered artifacts.

| Evidence set                      | Artifacts              | Partial frontiers | What it licenses                                                                                                                                               |
|-----------------------------------|------------------------|-------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Current registered $\alpha = .05$ | 7                      | 0                 | Exact per-candidate certificates; all are step-like.                                                                                                           |
| Exploratory or looser- $\alpha$   | 3                      | 2                 | Diagnostic evidence only; not the registered claim.                                                                                                            |
| NESTFUL refusal                   | 12 synthesized         | 0                 | $n_{\max} = 26 < 92$ ; no admission.                                                                                                                           |
| API-Bank refusal                  | 2 synthesized          | 0                 | $n_{\max} = 8 < 92$ ; no admission.                                                                                                                            |
| BFCL v4 refusal                   | 4 synthesized          | 0                 | $n_{\max} = 15 < 92$ ; no admission, pre-registered.                                                                                                           |
| AppWorld admission                | 1 admitted / 4 retired | 0                 | $n_{\max} = 136 \geq 92$ ; the one reachable external bound. Admitted gate is step-like; retirements are marginal ( $U = 0.051$ at 89 and 90 eligible groups). |

**Admissible is not dispatchable.** Table 5 separates the two. The admitted region is eligible on essentially every full-code trajectory and essentially no ReAct or plan-and-execute trajectory, because those architectures open by reading API documentation—3,752 of the 8,190 released trajectories begin with an API-description call—so the compiled prefix is no longer at position 0 and (2) refuses. A further 61 trajectories contain the program somewhere other than position 0, which the same invariant refuses. Full-code agents are eligible on 2,339/2,340 trajectories and iterative parallel function-calling agents on 762/1,170. Two secondary observations: the rate is flat across task difficulty (0.377 on `test_normal` against 0.380 on `test_challenge`), so this is a calling-discipline property rather than a workload property; and the maximal compilable depth is 2 on the eligible families, so the admitted two-step program sits at the ceiling the effect catalog permits rather than short of it. On the ineligible families the measured depth is a floor, not an estimate, because the API-documentation call is undeclared in the signed catalog and unknown effects are not reads.

#### F.4 WHAT THE REGISTERED GATE HAS AND HAS NOT DEMONSTRATED

This audit strengthens the evidentiary record precisely by preserving the negative result: the registered gate currently behaves as a support threshold, not as a demonstrated risk-coverage frontier. The exploratory partial frontiers therefore remain diagnostic rather than evidence for the main claim.




Figure 4: Per-family resource reductions behind table 1, plotted because the spread is the result: admissible prefix depth — not the optimizer — sets the savings, and the two deeper families sit at the feasibility ceiling of eq. (6).

## G EXTENDED LIMITATIONS

**A compression comparator ran, and did not engage.** Every headline resource number compares a compiled prefix against an *uncompressed* baseline, because deployed compression middleware reports token reductions of the same order on a comparable workload by an entirely different mechanism (Headroom Labs, 2026). A live Headroom v0.5.18 ablation adds `headroom_only/compiled_headroom` arms to both GitHub families on the existing 30-record held-out cohort (`paper/supplementary/headroom-ablation-protocol.md`). Headroom attempted every eligible JSON tool result and pre-model evidence object and applied *none* of them, saving zero tokens in either family, with all four paired quality contrasts holding 30/30 at exact  $p=1$ . This is a null engagement at this payload shape, not a measured composition or a measured floor — it does not establish Headroom’s behavior on larger contexts — but it does retire the narrower concern that our reduction is compression under another name: on this workload the two systems visibly do not touch the same bytes.

**A drift-parity protocol, pre-registered and not yet run.** The strongest standing objection is that a hand-written program reaches the same 90/90 contracts as the compiled artifact and is cheaper on one family. A three-arm design isolates the *induced verifier* from the *program* by pairing the compiled artifact with a permissive no-op verifier and testing both against the hand-written program under the nine-family metamorphic perturbation suite already implemented in `evaluation/perturb.py` (reordering, duplication, formatting, empty/null fields, schema drift, tool errors and timeouts), asking whether the compiled artifact abstains where an unverified program would answer silently (`paper/supplementary/drift-robustness-ablation-protocol.md`). Status: pre-registered on the deterministic demonstration substrate, not executed; no result from it is reported here.

**Where “unchanged baseline” is exact.** Pre-emission rejection is exact on both runtime paths; only a staging-owning runner can restore byte-identical model input after an attempt has been emitted (appendix B). Without one, a post-commit verifier failure is an incident rather than a fallback, which is why the compiled region is confined to the pre-commit prefix.

**How the calibration cohort must be sampled.** The calibration cohort must be sampled before filtering on whether the baseline happened to exhibit the mined region, with violation labels defined for every dispatch-eligible member; post-selecting only recurrent successes would not satisfy proposition 1. This is why the AppWorld substrate fixes its calibration share at exactly the 92 groups the bound requires and draws them by seeded shuffle over all clean tasks, rather than taking the tasks in which the mined family happens to have the deepest support.

Figure 4 plots the per-family spread behind table 1: admissible prefix depth, not the optimizer, sets the savings. The current empirical gate does not yet demonstrate a useful risk-coverage frontier at the registered 5% level. The strongest current runs are mostly all-or-none support tests (fig. 5),




Figure 5: NESTFUL family supports against the 92-group requirement implied by eq. (7) at  $\alpha = 0.05$ ,  $\delta = 0.1$ ,  $|\Lambda| = 11$ . The largest family reaches 26 groups, so every family retires at calibration even though provenance, synthesis, and replay all succeeded.

and the compiler-wide candidate-search problem is multiplicity-corrected only when the search is frozen (corollary 2); the primary families’ dominance search over two candidates is not. Likewise, exact task contracts do not certify full semantic equivalence of the final answer, especially once a downstream model continuation is involved. Two further operational caveats belong in the record: the headline live artifact was compiled without a sandbox available to the perturbation suite, so it records `perturbations_claimed: false` rather than a completed challenge, and registry signature verification was configured off in the reported run. These are not hidden defects in the repository; they are the reasons the paper avoids stronger deployment claims.

## H PROSPECTIVE GATE-FRONTIER PROTOCOL

paper/supplementary/prospective-gate-frontier-protocol.md pre-registers, but does not run, the experiment that could turn section 7’s step-like gate into a genuine risk-coverage frontier. **Status: PRE-REGISTERED, NOT EXECUTED.** No repository beyond what this paper already reports has been acquired under it, no provider call has been made, and no held-out case has been selected; this appendix summarizes the committed design so it cannot be revised in light of results it does not yet have.

**Cohort.** At least three repositories or domains, extending the provider-free `github_multirepo_preflight.py` harness, with at least 300 sealed held-out paired cases (zero compiled-only failures in 300 gives a one-sided 95% upper bound near 1%, well below the current small-cohort bounds), reported pooled and per-repository. **Compilation** freezes candidate selection throughout, so any admitted gate inherits corollary 2 rather than a per-candidate certificate. **Three arms** compared under identical records, model, cache policy, and ordering: the unchanged agent, the learned gate of section 2, and a support-only frequency gate. **At least three distinct nonzero coverage levels** must be observed for the learned gate, which requires a cohort deliberately including record classes expected to sit at different points on a certification-difficulty gradient, not two endpoints from a homogeneous task. Thresholds, features, splits, hypotheses, and the stopping rule are fixed and hashed before acquisition.

**Decision rule.** A frontier with wrong-dispatch at or below the registered bound at every level, dominating the support-only gate at matched coverage, is the predicted positive outcome. A gate that stays step-like while the support-only comparator does not (or vice versa) is a reported negative for the learned score specifically. Any held-out wrong dispatch on either gate is an adverse finding reported in full, ahead of every other reading. Two extensions are named and explicitly deferred pending their own feasibility bar: model/provider breadth, and an executable workflow-compiler comparator (AWO is the closest published candidate), run only if it can match records, model, execution placement, overhead accounting, cache policy, and grading without hand-waving — a forced comparator would weaken the paper more than omitting one.